


Introduction

Computer Engineering

Hardware

- CPU Central Processing Unit
- Memory Stores instructions and data
- Input / Output Interface to external devices
- System-Bus Electrical connection of blocks



Datapath

- ALU
- Registers

Arithmetic and Logic Unit Fast but limited storage inside CPU

Control Unit


- Finite State Machine

Reads and executes instructions

- Types of instructions Data transfer, Arithmetic, logical and jumps

Software

From C to executable



1. Preprocessor
 - Text processing
 - Pasting of #include files
 - Replacing macros (#define)
2. Compiler
 - Translate CPU-independent C-code into CPU-specific assembly code
3. Assembler
 - Translate to machine instructions
 - Result: Relocatable object file
 - Binary file → not readable with text editor
4. Linker
 - Merge object files
 - Resolve dependencies and cross-references
 - Create executable

Cortex-M Architecture

Registers

- 16 Core Registers
- 32-Bit wide

- R0-R7 Lower Registers
- R8 - R12 Higher Registers
- R13 Stack Pointer Temp Storage
- R14 Link Register Return from Procs
- R15 Program Counter Addr of next Instr.

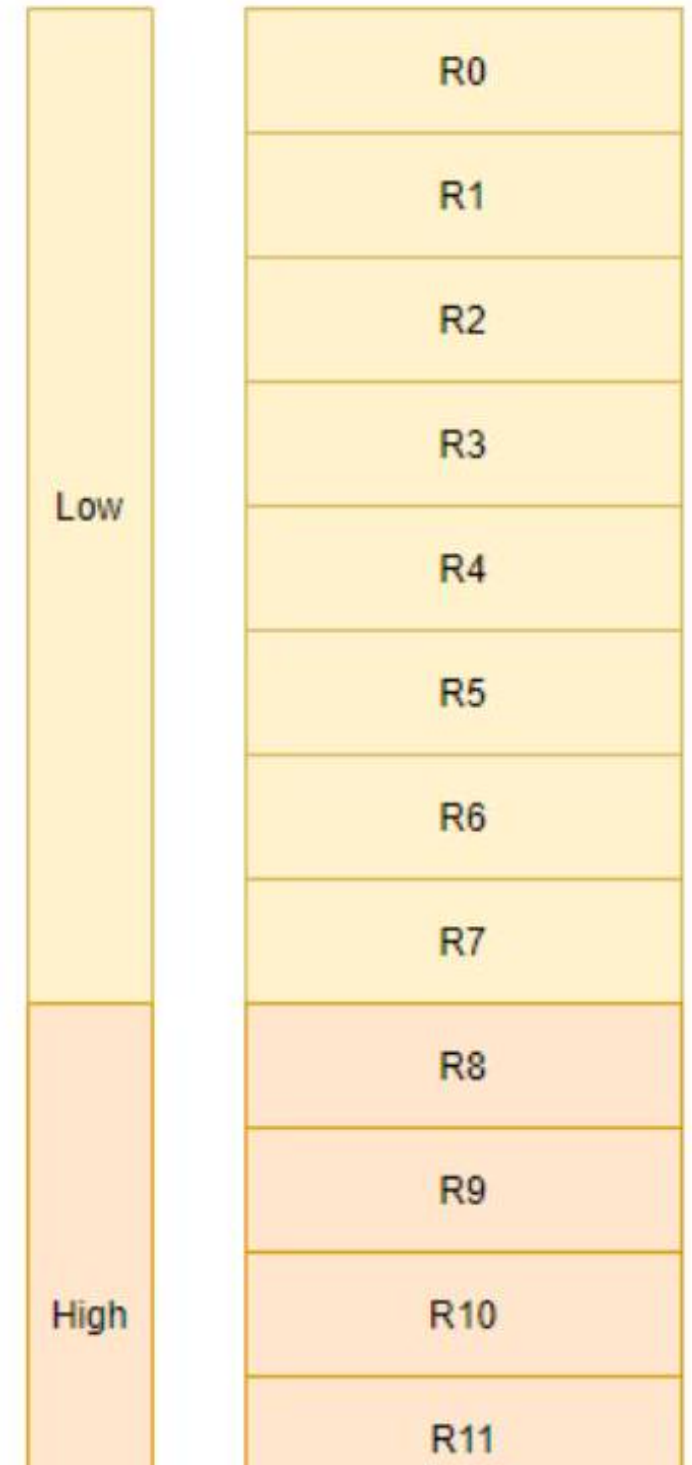
ALU

- 32-Bit wide processing unit

APSR (Flag Register)

- N Negative
- Z Zero
- C Carry
- V Overflow

Instruction Set



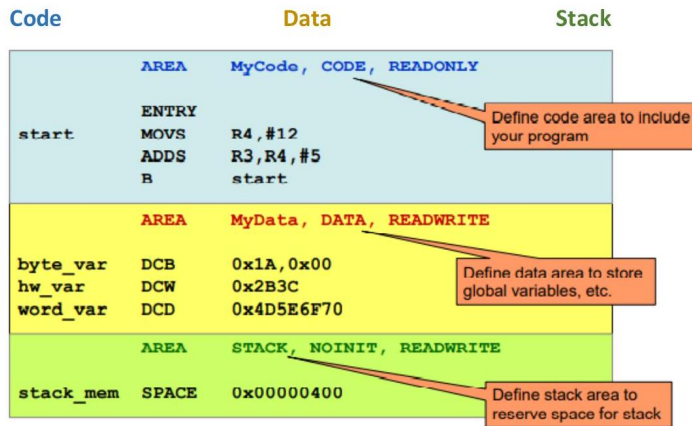
- 16-Bit Thumb instruction encoding

Label	Instr.	Operands	Comments
demoprg	MOVS	R0,#0xA5	; copy 0xA5 into register R0
	MOVS	R1,#0x11	; copy 0x11 into register R1
	ADDS	R0,R0,R1	; add contents of R0 and R1

Instruction Types

- Data transfer
- Data processing
- Control flow Move, Load and Store Arithmetic, Logical and Shift operations Branches and functions

Assembly Program Structure



Directives for initialized data

- DCB Bytes
- DCW Half-Words
- DCD Words

var1	DCB	0x1A			
var2	DCB	0x2B	0x3C	0x4D	0x5E
var3	DCW		0x6F70	0x8192	
var4	DCD		0xA3B4C5D6		

Directives for uninitialized data

- SPACE Bytes to be reserved

Data Transfer

Data Transfer Instructions

Loading Data

- MOVS
- Reg to Reg MOVS R1, R2
- 8-Bit Literal MOVS R1, #0x1C
- Constant MOVS R1, #MyConst
- LDR
- 32-Bit Literal LDR R1, #0xA1B2C3D4
- Literal + Offset LDR R1, [PC, #12]
- Constant LDR R1, =MyConst
- Reg Value LDR R1, [R2]

- LDRB
- Load Register Byte
- Bits 31 to 8 set to zero
- LDRH
- Load Register Half-word
- Bits 31 to 16 set to zero

Load Array

- my_array = 3 * 4 Bytes
- Instructions = 5 * 2 Bytes
- Literals (0x08) = 1 * 4 Bytes

```

AREA my_data, DATA, READWRITE
00000000 11223344 my_array DCD 0x11223344
00000004 55667788 DCD 0x55667788
00000008 99AABBCD DCD 0x99AABBCD

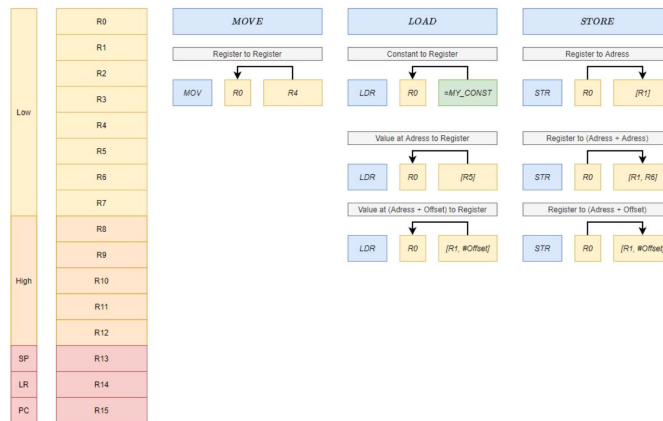
```

```

AREA myCode, CODE, READONLY
. . .
0000007C 4906 LDR R1, my_array ; load address of array
0000007E 4B07 LDR R3, #0x08

; indirect addressing
00000080 680C LDR R4, [R1] ; base R1
00000082 684D LDR R5, [R1, #0x04] ; base R1, immediate offset
00000084 58CE LDR R6, [R1, R3] ; base R1, offset R3

```



Storing Data

- STR
- Value from Register STR R1, [R2]
- Value from Reg + Offset STR R1, [R2, #0x04]
- STRB
- Store Register Byte (Low 8 bits of register stored)
- STRH
- Store Register Half-word (Low 15 bits of register stored)

Arithmetic Operations

Arithmetic Operations

Flags (APSR = N, Z, C, V)

Instructions ending with with «S» allow flag modification

- ADDS
- SUBS
- MOVS
- LSLS

Flag	Meaning	Action	Operands
Negative	MSB = 1	N = 1	signed
Zero	Result = 0	Z = 1	signed, unsigned
Carry	Carry	C = 1	unsigned
Overflow	Overflow	V = 1	signed

Overview

- ADD / ADDS
- ADCS Addition with Carry
- ADR Address to Register
- SUB / SUBS
- SBCS
- RSBS
- MULS

Subtraction

Subtraction with carry (borrow)

Reverse Subtract (negative)

Multiplication

$A + B$

$A + B + c$

$PC + A$

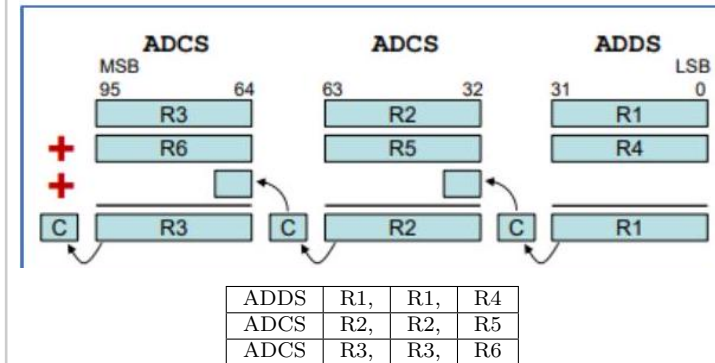
$A - B$

$A - B - !c$

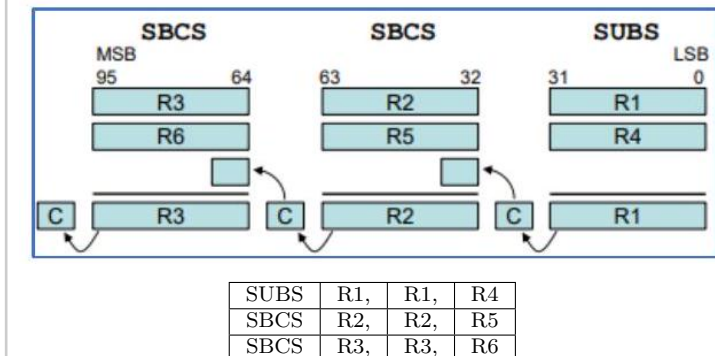
$-1 \cdot A$

$A \cdot B$

Multi-Word Addition with ADCS



Multi-Word Subtraction with SBCS



Negative Number

- 2' Complement $A = !A + 1$

Carry and Overflow

unsigned

- Addition → $C = 1 \rightarrow$ carry result too large for available bits
 - Subtraction → $C = 0 \rightarrow$ borrow result less than zero → no negative numbers
- signed
- Addition → potential overflow in case of operands with equal signs
 - Subtraction → potential overflow in case of operands with opposite signs

Addition and Subtraction


- Addition $C = 1 \rightarrow$ Carry

1	1	0	1	13d
0	1	1	1	7d
1	1	1	1	
1				
1	0	1	0	0

- Subtraction $C = 0 \rightarrow$ Borrow sign
 $6d - 14d = 0110b - 1110b = 0110b + 0010b$

0 1 1 0 6d
0 0 1 0 2d = TC(14d)

0100



Logic, Shift and Rotate Instructions

Logical Instructions

The following instruction only affect N and Z flags

- ANDS

- BICS
- EORS
- MVNS Bitwise NOT
- ORRS

Bitwise OR

Rdn # Rm

Rdn & Rm

Rdn & ! Rm $a \& \sim b$


Rdn \$ Rm $a \wedge b$

Rm a

a | b

Shift Instructions

- LSLS Logical Shift Left $2^n \cdot Rn \ 0 \rightarrow LSB$
- LSRS Logical Shift Right $2^{-n} \cdot Rn \ 0 \rightarrow MSB$
- ASRS Arithmetic Shift Right $R^{-n} \pm MSB \rightarrow MSB$
- RORS Rotate Right $LSB \rightarrow MSB$



Sign-Extension

Add additional bits

- Unsigned zero extension fill left bits with zero
- Signed sign extension copy sign bit to the left

Unsigned → ~ Zero Extension			
1011 →	00001011	0011 →	00000011
Signed → Sign Extension			
1011 →	11111011	0011 →	00000011

Truncation

Cast cuts out the left most digits

- Signed possible change of sign
- Unsigned results in module operation

Integer ranges based on word sizes

8-bit	hex	unsigned	signed	16-bit	hex	unsigned	signed
	0 × 00				0 × 0000		
	0x75	127	127		0x7FFF	F 32'767	-1
	0x80	128	-128		0x8000	032'768	-32'768

	0xFF	255	-1		0xFFFF	F 65'535	-32'768
32-bit		hex	unsigned	signed	hex	unsigned	signed
		0x0000 0000	0	0			
		0x7FFF'FFFF	2'147'483'647	...			
				2'147'483'647			
		0x8000'0000	2'147'483'648 -	-2'147'483'648			
		0xFFFF'FFFF	4'294'967'295	...			
				-1			

Control Structures

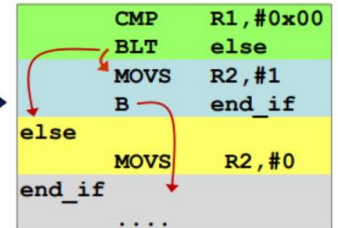
Branch Instructions

Unconditional Branches	Conditional Branches	Overview
<ul style="list-style-type: none">B (immediate)• Direct• Relative• BX (Branch and Exchange)• Branch and Exchange• Indirect• Absolute	<ul style="list-style-type: none">Flag-dependent and arithmetic branches• Indirect• Absolute	<p>Type</p> <ul style="list-style-type: none">Unconditional Branch alwaysConditional Branch if condition is met <p>Address hand-over</p> <ul style="list-style-type: none">Direct Target addresses part of instructionIndirect Target address in register <p>Address of target</p> <ul style="list-style-type: none">Relative Target address relative to PCAbsolute Absolute address <p>Compare and Test</p> <ul style="list-style-type: none">TST AND without changing the valueCMP SUBS without changing the value

Selection (IF-ELSE)

```
int32_t nr;  
int32_t isPositive;  
...  
if (nr >= 0) {  
    isPositive = 1;  
}  
else {  
    isPositive = 0;  
}
```

Assume: nr in R1
isPositive in R2



Switch

Jump Table

```
uint32_t result;  
case 0: -1  
case 1:  
    result += 17; break;  
    result += 13; //fall through  
case 3: case 5: result += 37; break;  
default:  
    result = 0;  
}
```

NR_CASES case_switch	EQU	6
	CMP	R1, #NR_CASES
	BHS	case_default
	LSLS	R1, #2 ; * 4
	LDR	R7, =jump_table
	LDR	R7, [R7, R1]
	BX	R7
case_0	ADDS	R2, R2, #17
case_1	ADDS	R2, R2, #13
case_3_5	ADDS	R2, R2, #37
	B	end_sw_case
case_default	Movs	R2, #0
end_sw_case	...	
jump_table	DCD	case_0
	DCD	case_1
	DCD	case_default
	DCD	case_3_5
	DCD	case_default
	DCD	case_3_5

C	Assembly
#include <utils_ctboard.h> #include <stdint.h> ... int32_t i = 0; int32_t count = 0; for(i = 0; i < 10; i++) { count++; }	AREA progCode, CODE, READONLY THUMB PROC EXPORT main main LDR R6, #i ; R6=address of i LDR R0, [R6] ; R0=value at i LDR R7, =count ; R7=address of count LDR R1, [R7] ; R1=value at count B cond loop ADDS R0, R0, #1 ADDS R1, R1, #1 cond CMP R0, #10 ; *signed* comparison BLT loop ; store final i STR R0, [R6] ; store final i STR R1, [R7] ; store final count endless B endless ENDP AREA progData, DATA, READWRITE i DCD 0 count DCD 0 END

Subroutines and the Stack

Subroutine Call and Return

- Label with Name (MulBy3)
- Return Statement (BX LR)

00000050	4604	MulBy3	MOV	R4, R0
00000052	0040	LSLS	RO, #1	
00000054	4420	ADD	R0, R4	
00000056	4770	BX	LR	

Stack

- Stack Area (Section)
- Stack Pointer (SP)
- PUSH {...}
- POP {...}
- Direction on ARM
- Alignment
- Only words

Continuous area of RAM

R13 → points to last written data value

Decrement SP and store words

Read words and increment SP

full-descending stack

word-aligned

32-Bit

Stack - Push and Pop

- Number of Pushes = Number of Pops
- Stack-limit < SP < stack-base

ADDR_LED_31_0	EQU	0x60000100
LED_PATTERN	EQU	0xA55A5AA5
subrExample	PUSH	{R4, R5, LR}
; write pattern to LEDs		
LDR	R4, =ADDR_LED_31_0	
LDR	R5, =LED_PATTERN	
STR	R5, [R4]	
BL	write7seg	
POP	{R4, R5, PC}	

Save LR and registers used by subroutine

Call another subroutine

Restore registers and PC

PUSH {R2,R3,R6}

00000000	B083	SUB	SP, SP, #12
00000002	9200	STR	R2, [SP]
00000004	9301	STR	R3, [SP, #4]
00000006	9602	STR	R6, [SP, #8]

POP {R2,R3,R6}

00000008	9A00	LDR	R2, [SP]
0000000A	9B01	LDR	R3, [SP, #4]
0000000C	9B02	LDR	R6, [SP, #8]
0000000E	B003	ADD	SP, SP, #12

Parameter Passing

Where

- Register
- Global variables
- Stack
- Caller: PUSH parameter on stack
- Callee: Access parameter through LDR

Reentrancy

- Recursive Function Calls
- Registers and global variables are overwritten
- Requires an own set of data for each call
- Solution:
- ARM Procedure Call Standard

Passing through global variables

- Shared variables in data area
- Overhead to access variable
- Error prone, unmaintainable
- By reference
- Allows passing of larger structures

ARM Procedure call Standard

Parameters

- Caller copies arguments From R0 to R3
- Caller copies additional parameters to stack

Returning fundamental data types

- Smaller than word
- Word
- Double word
- 128-Bit
- zero or sign extend to word return in R0 / R1 return in R0 - R3

Returning composite data types

- Up to 4 bytes return in R0
- Larger than 4 bytes stored in data area

Loops

- Do while: Post-Test Loop

```
int32_t nr;
int32_t sum;
...
sum = 0;
do {
    sum += nr;
} while (sum < 100);
```

Assume: nr in R1
sum in R2


```
MOVS R2, #0
loop ADDS R2, R2, R1
    CMP R2, #100
    BLT loop
....
```

- While = Pre-Test Loop

```
int32_t nr;
int32_t prod;
...
prod = 1;
while (prod < 100) {
    prod *= nr;
}
```

Assume: nr in R1
prod in R2

```
MOVS R2, #1
B test
loop Muls R2, R1, R2
    test
    CMP R2, #100
    BLT loop
....
```



- For = Pre-Test Loop

Register / "pass by value"

```
AREA exData, DATA, ...
```

```
...
```

```
AREA exCode, CODE, ...
```

```
...
```

```
MOVVS R1, #0x03
```

```
BL double
```

```
MOVVS ..., R0
```

```
...
```

```
double
```

```
LSLS R0, R1, #1
```

```
BX LR
```

caller

callee
function
double

Register Usage


Register	Synonym	Role
r0	a1	Argument / result / scratch register
r1	a2	Argument / result / scratch register
r2	a3	Argument / scratch register
r3	a4	Argument / scratch register
r4	v1	Variable register 1
r5	v2	Variable register 2
r6	v3	Variable register 3
r7	v4	Variable register 4
r8	v5	Variable register 5
r9	v6	Variable register 6
r10	v7	Variable register 7
r11	v8	Variable register 8
r12	IP	Intra-Procedure-call scratch register
r13	SP	
r14	LR	

Register contents
might be modified
might be modified

Callee must preserve contents
of these registers (Callee saved)

Subroutine Call – Caller Side

Pattern as used by the compiler. Manually written assembly code may be slightly different.



On return from subroutine

Modular Coding and Linking

Modular Coding / Linking

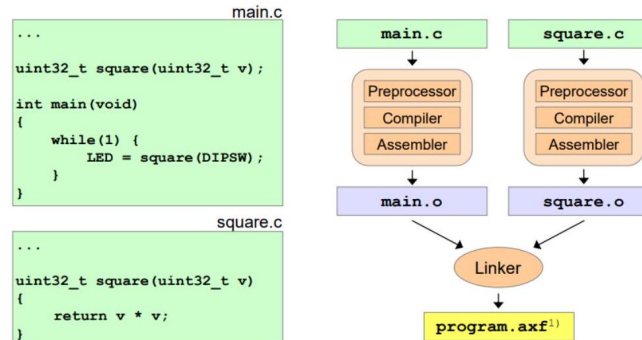
From source code to executable program

Compile / assemble each module

- Results in an object file for each module

Link all object files

- Results in one executable file



Managing complexity by modular programming

Topic	Benefits
Enable working in teams	Multiple developers working on the same code repository
Useful partitioning and structuring of the programs	Eases reusing of modules
Individual verification of each module	Benefits all users of the module
Providing libraries of types and functions	For reuse instead of reinvention
Mixing of modules that are programmed in various languages	E.g. mix C and assembly language modules
Only compile the changed modules	Speeds up compilation time

ARM assembly IMPORT and EXPORT keywords

Linkage control

- EXPORT for use by other module
- IMPORT from another module for use in this module Internal symbols
- Neither IMPORT nor EXPORT

usable outside of module main

```

; main.s
AREA myCode, CODE, READONLY
EXPORT main
IMPORT square
main PROC
    LDR    r0, a_addr
    LDR    r0, [r0, #0] ; a
    BL     square
    ...
ENDP
a_addr DCD    a
b_addr DCD    b

AREA myData, DATA
a DCD    0x00000005
b DCD    0x00000007

```

from module square

Linker Input - Object files

Code section Code and constant data of the module, base at address 0 × 0
Data section All global variables of the module, based at address 0 × 0
Symbol table All symbols with their attributes like global/local, reference
Relocation table

Which bits of the data and code section need to be adjusted (and how) after merging the sections in the linking process

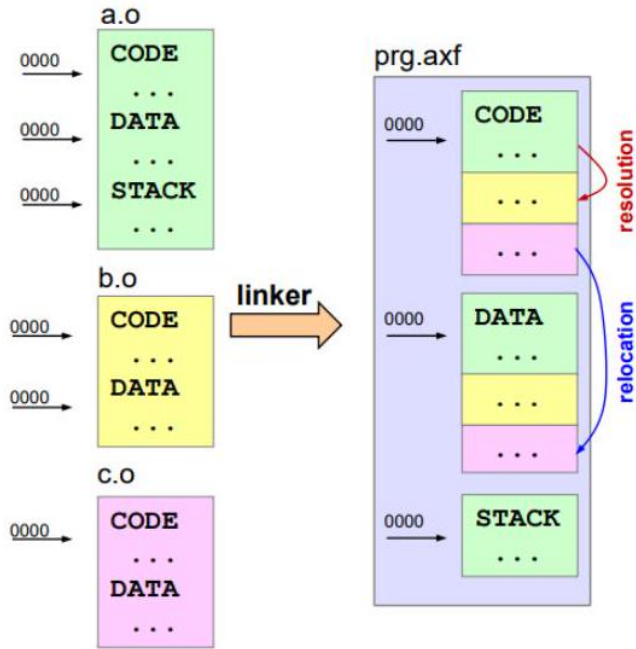
ARM tool chain uses ELF for object files

Linker tasks

- Merge object file code sections
- Merge object file data sections
- Symbol resolution
- Address relocation

Linker Output

- AXF = **ARM** eXecutable File



Exceptional Control Flow

Interrupt sources

- Peripherals signal to CPU that an event needs immediate attention
- Can alternatively be generated by software request
- Asynchronous to instruction execution

System exceptions

- Reset
- NMI
- Faults
- System Level Calls

Restart of processor

Non-maskable Interrupt (cannot be ignored) Undefined instructions OS calls - Instructions SVC and PendSV

PRIMASK

- Single bit controlling all maskable interrupts

- Disable	set PRIMASK	clear PRIMASK	Assembly
- Enable			CPSID ¹⁾ CPSIE ¹⁾

On reset PRIMASK = 0 → enabled

Storing the context

Interrupt event can take place at any time

- E.g. between TST and BEQ instructions
- ISR call requires automatic save off regs and caller saved registers

ISR call

- Stores xPSR, PC, LR, R12, R0-R3 on Stack
- Stores EXC_RETURN to LR

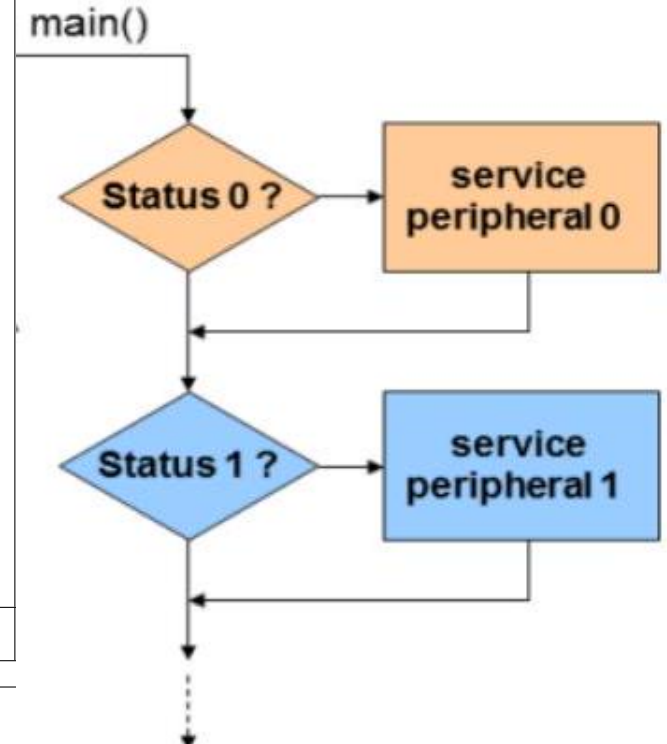
ISR Return

- Use BX LR or POP {..., PC}
- Loading EXC Return into PC
- Restores R0-R3, R12, LR, PC and xPSR from Stack

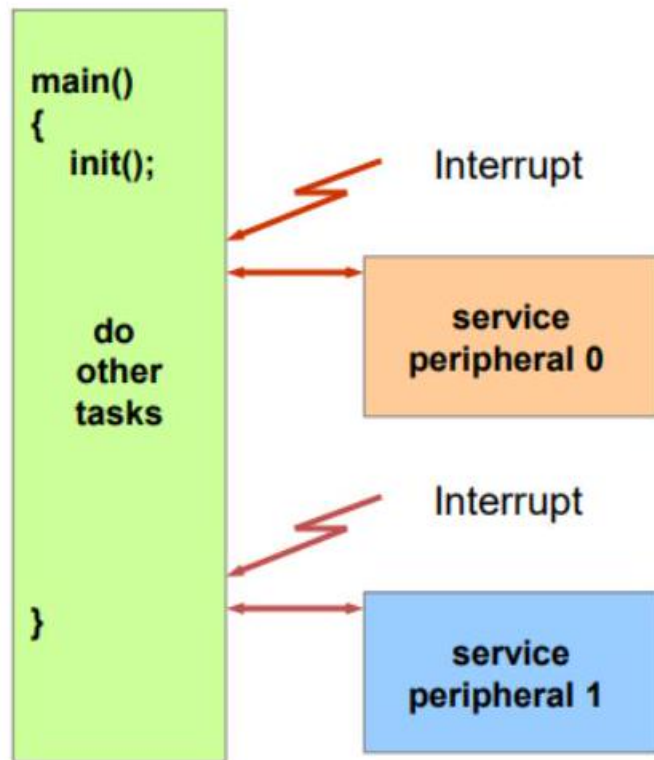
Polling

Periodic query of status information

- Reading of status registers in loop
- Synchronous with main program
- Advantages
- Simple straightforward
- Implicit synchronisation
- Deterministic



- No additional interrupt logic required
- Disadvantages
- Busy wait -> wastes CPU time
- Reduced throughput
- Long reaction time
- No synchronization
- Difficult debugging



RISC = Reduced Instruction Set Computer

- Few instructions, unique instruction format
- Fast decoding, simple addressing
- Less hardware -> allows higher clock rates
- More chip space for registers (up to 256!)
- Load-store architecture reduces memory access,

CPU works at full-speed on registers

- Higher clock frequencies
- Easy and shorter pipelines (instruction size / duration)

RISC

- Load / Store Architecture
- Data processing instructions only available on registers

Example: **Balance = Balance + Credit**

```

LDR R0,=Credit
LDR R1,[R0]
LDR R0,=Balance
LDR R3,[R0]
ADDS R2,R1,R3
STR R2,[R0]

```

CISC

- One of the operands of an instruction may directly be a memory location

```

MOV AX, [Credit]
ADD [Balance], AX


```

CISC = Complex Instruction Set Computer

- More complex and more instructions
- Less program memory needed with complex instructions
- Short programs may work faster with less memory accesses


Von Neuman Architecture

- Same memory holds program and data
- Single bus system between CPU and memory Systembus

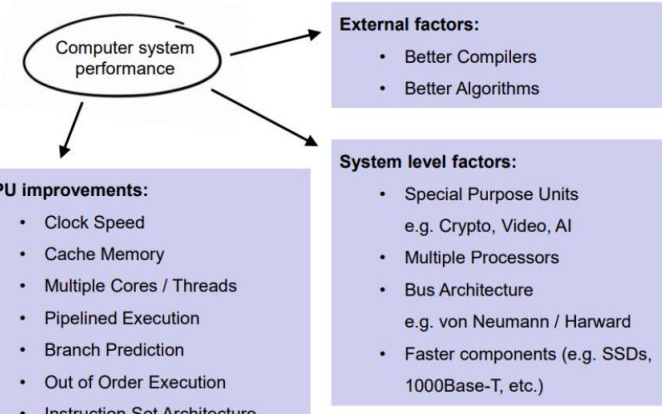


Harvard Architecture


- «Mark I» at Harvard University
- Separate memories for program and data
- Two sets of addresses/data buses between CPU and memory



How to Increase System Speed?

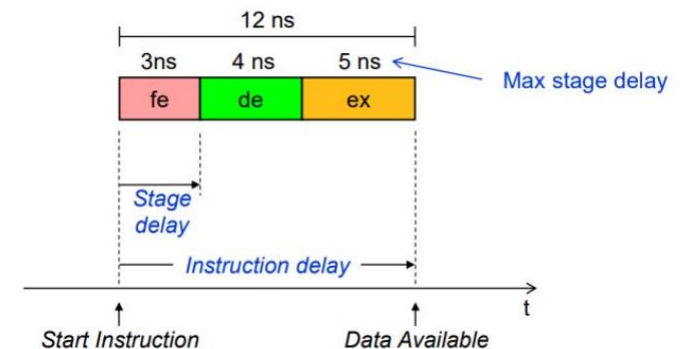


Fetching the next instruction, while the current one decodes
Sequential vs. Pipelined



Timings and definitions (Example)

- Fe: fetch Read instructions 3 ns
- De: decode Decode instruction, read register or memory 4 ns
- Ex: execute Execute instruction, write back result 5 ns



Advantages of pipelining

- All stages are set to the same execution time
- Massive performance gain
- Simpler hardware at each stage allows for a higher clock rate

Disadvantages

- A blocking stage blocks while pipeline
- Multiple stages may need to have access to the memory at the same time

Instructions per second

Without pipelining

$$\frac{\text{Instructions}}{\text{second}} = \frac{1}{\text{Instruction delay}}$$

With pipelining

- Pipeline needs to be filled first

Increasing System Performance

Speed vs Low Power Aspects of Optimization	
Optimizing for Higher speed	Drawbacks on Power, cost, chip area
Lower cost	Speed, reliability
Zero power consumption	Speed, cost
Super reliable	Chip area, cost, speed
Temperature range	Power, cost lifetime

- After filling, instructions are executed after every stage

$$\frac{\text{Instructions}}{\text{second}} = \frac{1}{\text{Max stage delay}}$$

Optimal pipelining

- All operations here are on registers
- In this example it takes 6 clock cycles to execute 6 instructions
- Clock cycles per instruction (CPI) = 1

Cycle		1	2	3	4	5	6	7	8	9
Operation										
ADD		fe	de	ex						
SUB			fe	de	ex					
ORR				fe	de	ex				
AND					fe	de	ex			
ORR						fe	de	ex		
EOR							fe	de	ex	

Special situation: LDR

- In this example it takes 7 clock cycles to execute 6 instructions
- Read cycle must complete on the bus before LDR instruction can complete
- Next 2 instructions must wait one pipeline cycle (S = stall)
- Clock cycles per Instruction (CPI) = 1.2


Control Hazards

- Branch / jump decisions occur in stage 3 (ex)
- Worst case scenario - conditional branch taken:

```

00000000 4288      CMP    R0,R1    ; R0 > R1 ?
00000002 D802      BHI     go_on    ; if higher -> go_on
00000004 000A      MOVS    R2,R1    ; otherwise exchange regs
00000006 0001      MOVS    R1,R0
00000008 0010      MOVS    R0,R2
0000000A 2305      MOVS    R3,#5
0000000C ...

```



Reduce control hazards

- Loop fusion reduces control hazards

Ideas to further improve pipelining

- Branch prediction
- Store last decisions made for each conditional branch
- -> probability is high that the same decision is taken again
- Instruction prefetch
- Fetch several instructions in advance
- -> better use of system bus
- -> possibility of «Out of Order Execution»
- Out of Order Execution
- If one instruction stalls, it might be possible to already execute the next instruction

Limits of optimization

- Complex optimizations -> sever security problems

- Instructions executed, that would throw access violations under «In Order» circumstances.
- «Meltdown» and «Spectre» attacks: allow a process to access the data of another process

Parallel Computing

- Streaming / Vector processing One instruction processes multiple data items simultaneously
- Multithreading Multiple programs/threads share a single CPU
- Multicore Processors One processor contains multiple CPU cores
- Multiprocessor Systems A computer system contains multiple processors

