



Looking for People

- Safe and Dependable Machine Learning
 - FPGA and SW
- Safe Vision
 - FPGA and SW
- Drone/Autonomous Vehicles
 - SW / HW

- Looking for
 - 1-2 full time HW/SW and/or 1-2 MSE SW

IoT Lecture 7



<https://www.cisoplatform.com/profiles/blogs/classification-of-iot-devices>
<https://fr.wikipedia.org/wiki/Ethernet>
https://en.wikipedia.org/wiki/Protocol_stack

Constrained Devices – Agenda

- RFC definition of constrained device
 - Benchmarking performance on constrained devices
- Ethernet Communication
 - Physical layers – Fast Ethernet and Advanced Physical Layer
 - HW footprint
 - Serial to parallel conversion
 - Data transfer to memory
 - Data buffering schemes
- Communication stacks
 - Zero copy stack
 - Linux IP stack
 - Socket interfaces

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power


Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

IoT Lecture 7 – Constrained Devices



<https://www.cisoplatform.com/profiles/blogs/classification-of-iot-devices>

Constrained Device Definitions (1)

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack

- RFC 7228 defines terminology of constrained devices
 - Constrained -> resource constrained
- Cost of devices silicon cost based on
 - Size of wafer, area of die, yield after wafer production
 - Number of pins, yield after packaging

	Name	data size (e.g., RAM)	code size (e.g., Flash)
Not IP Capable (?)	Class 0, C0	<< 10 KiB	<< 100 KiB
CoAP Capable	Class 1, C1	~ 10 KiB	~ 100 KiB
	Class 2, C2	~ 50 KiB	~ 250 KiB

Table 1: Classes of Constrained Devices (KiB = 1024 bytes)

Constrained Device Definitions (2)

■ Also categorized devices by energy/power requirements

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack

Name	Definition	SI Unit
P_s	Sustainable average power available for the device over the time it is functioning	W (Watt)
E_t	Total electrical energy available before the energy source is exhausted	J (Joule)

Table 2: Quantities Relevant to Power and Energy

Name	Type of energy limitation	Example Power Source
E_0	Event energy-limited	Event-based harvesting
E_1	Period energy-limited	Battery that is periodically recharged or replaced
E_2	Lifetime energy-limited	Non-replaceable primary battery
E_9	No direct quantitative limitations to available energy	Mains-powered

Table 3: Classes of Energy Limitation

Name	Strategy	Ability to communicate
P_0	Normally-off	Reattach when required
P_1	Low-power	Appears connected, perhaps with high latency
P_9	Always-on	Always connected

Table 4: Strategies of Using Power for Communication

Constrained Device Experience (1)

■ Bare Metal

- Bit-Bang serial UART, 8085, 4MHz., bare metal, assembler,
 - 32k code EPROM (ca. 2k used), 128 byte RAM used
- CANopen devices, 20+ MHz., bare metal
 - 4-5k code -> 25 k -50 k code.
 - ~100-200 Bytes data -> 0-5k – 1k data
- Ethernet Powerlink, 75MHz. ARM, bare metal (no IP stack)
 - ~250k code, ~100k data
- General expectation that IP stack requires OS (multi-tasking)

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

Constrained Device Experience (1)

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack

- Typical embedded OS
 - eCOS – 80k kernel
 - BSD2000 TCP/IP stack – 350 kByte
- TCP/IP stack is fundamentally large – problem for constrained devices
 - Several lite IP-Stacks in the embedded market
 - (Coldfire 45 kBytes code as a webserver and 16 kBytes stack+heap)
 - Also HW IP, UDP stacks available
- Provisioning
 - What HW capabilities are required to
 - Provide required performance
 - “future-proof” the application


Provisioning and Performance (1)

■ Provisioning


@16 ms
cycle time
34% idle

■ ARM9 666MHz. eCOS - Altera

- 18% idle time @666MHz for 1ms?
 - WTH



@1 ms
cycle time
18% idle



Constrained Devices – Performance (2)

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack

■ Macro and Micro-Benchmarking

- Determine suitability of platform for a task

■ Micro-Benchmarks

- Used to determine general platform capabilities
- Lmbench can be used to test (Unix-like) platforms
- Processor and Processes
- Integer / floating point operations
- Local communication latencies and bandwidths
- File and VM latencies
- Memory Latencies


Constrained Devices – Performance (3)

■ For instance

- Raspberry PI versus BeagleBone
- Different system variations in idle mode, memory usage and CPU usage
- SD4 and SD10 -> secure SD card for bootloading and filesystem

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack



Constrained Devices – Performance (4)

■ Macro-Benchmarking

- Application orientated benchmarking


Trigger on HW signal generated by write to output – first line of DSR (Driver under eCOS)

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack

Trigger on Interrupt – HW signal generated by input buffer in FPGA. 30'000 Frames

Trigger on HW signal generated by write to output – first line of ISR



Constrained Devices – Performance (5)


■ Loop- Back Times, UDP frames, two platforms

- Loopback done by driver (DSR)
- Similar results for socket-level loopback




Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer
- Session 3: Communication Stacks
- Communication Stack

Xilinx A9 666 MHz. eCOS
30'000 frames



NIOS II 100MHz. eCOS.
30'000 frames



IP benchmarking

- RFC 2544 – IP independent benchmarking
- RFC 5180 – IPv6 benchmarking specifica

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power


Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

IoT Lecture 7 – Ethernet Communication



<https://fr.wikipedia.org/wiki/Ethernet>

Physical layers (1)

- Already met CAN and RSxxx
 - Fieldbus physical layers
- Ethernet
 - IEEE 802.3 – 10 MBit/s – largely obsolete
 - IEEE 802.3u – Fast Ethernet – 100 MBit/s – in industry
 - IEEE 802.3ab – Gigabit Ethernet – 1 Gbit/s – in office
 - IEEE 802.3cg – Advanced Physical Layer – in industry
- In terms of constrained devices, Ethernet and IoT
 - Physical footprint
 - Physical size AND power requirements of the integrated circuit (IC) and support circuitry
 - Maximum rate of data consumption by microcontroller
 - Ability to shift frames into memory
 - Ability to process frame headers
 - Ability to process frame data

Learning Aim: The student will know the four main Ethernet types in IoT/industry use and their IEEE standards numbers

The Layer System

Session 1: Constrained Devices


- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication**
- Physical Layer
- Session 3: Communication Stacks**
- Communication Stack

■ OSI layer model

- Used to understand communication systems
- Difficult to program in purity

■ Ethernet Relevance

- Application Layer -> Application hands over data that needs to be sent
- Presentation -> how the data is to be formatted for transfer to application
- Session -> the context of the data transfer
- Transport -> the connection
- Network -> the logical reachability
- Data Link -> the physical connection
- Physical layer -> assemble bits to bytes



Learning Aim: The student will understand the layering system

<http://www.upnp.org/specs/arch/UPnP-arch-DeviceArchitecture-v1.1.pdf>

Footprint - Connector

Connector

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication



- Physical Layer

Session 3: Communication Stacks

- Communication Stack



**M12 4-pin D-Coded
(M12 8-pin A-Coded)**



10Gbit/s



Automotive connector

Learning Aim: The students will be able to describe the practical issues around Ethernet connector technologies

Footprint - Filter / Magnetics

- Differential signal input (no DC element)
- RC filters
- Transformers for
 - Galvanic isolation (no common ground)
 - Common mode rejection (robustness to noise)

Session 1: Constrained Devices



- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack




Learning Aim: The students will know the purpose of the filter/magnetics circuits

http://electronicdesign.com/site-files/electronicdesign.com/files/uploads/2014/04/0614_EnergyZarr_FO-big.gif

The PHY (1)

- PHYSical Layer Device
- Mixed signal device
- Converts serial bitstream of 5B @125MHz to nibble stream @ 25 MHz.
 - 5B @ 125 MHz → 4B @ 100MHz.
 - 4B serial @ 100MHz -> 4 Bit parallel @ 25MHz



Learning Aim: The student will know the purpose of the PHY device

Learning Aim: The student will be able to explain the frequencies from wire to MAC

<http://donar.messe.de/exhibitor/hannovermesse/2015/E894782/gallery-868x0-301096.jpg>

The PHY (2)

Session 1: Constrained Devices


- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication**
- Physical Layer
- Session 3: Communication Stacks**
- Communication Stack

- Serialiser, scrambler, NRZI converter and MTL3 encoder
- Usually Digital Signal Processor (DSP) as sampling -> decoding element
- Due to DSP, modern PHY's can handle over 100m (f.i. 125 meters) cable lengths

Data interface to MAC

- Media Independent Interface (MII)
- TX_CLK (from PHY)
- TX_D [0 .. 3] -> 25MHz out
- TX_EN – enable transmission (MAC to PHY)
- TX_ER – transmit error (MAC to PHY – kill packet)
- RX-CLK (from PHY)
- RX_D [0 .. 3]
- RX_DV – receive data valid
- RX_ER – Receiver error
- CRS – Carrier sense
- COL – Collision Detect
- **Problem:** MII has high pin count -> Reduced Media Independent Interface (RMII)

Control interface between PHY and MAC/uP



- Phase-Locked Loop generates sampling frequency
 - Also generates heat
 - PHY usually as separate component

Learning Aim: The student will know the general components of the PHY device and its interfaces

The PHY (5)


- Session 1: Constrained Devices
 - Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
 - Physical Layer
- Session 3: Communication Stacks
 - Communication Stack

■ GMII

- GTX-CLK (from PHY)
- TX-CLK
- TX_D [0 .. 7] -> **25MHz out**
- TX_EN – enable transmission (MAC to PHY)
- TX_ER – transmit error (MAC to PHY)

- RX-CLK (from PHY)
- RX_D [0 .. 7]
- RX_DV – receive data valid
- RX_ER – Receiver error
- CRS – Carrier sense
- COL – Collision Detect

Takeaway: The physical footprint of the interface gets larger with increased transmission speed/bit-rate.



■ SGMII

- High speed serial interface differential DDR
- 625 MHz.
- GMII signals are encoded, serialised and transmitted with the data


http://www.angelfire.com/electronic2/sharads/protocols/MII_Proocols/sgmii.pdf

Advanced Physical Layer

- Session 1: Constrained Devices
 - Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
 - Physical Layer
- Session 3: Communication Stacks
 - Communication Stack

■ Beginnings: Automotive BroadReach

- 2-wire unshielded twisted pair
- 100Mbit/s full duplex
- Max 15 meters
- 4B/3B coding – aggressive filtering



■ Takeover: Process industry

- 10BASE-T1L IEEE 802.3cg-2019
- 2-wire unshielded pair
- Power-over-Ethernet (-like)
 - For EX-environments
- 10 Mbit/s full duplex
- Trunk and Spur architecture
 - Unpowered trunk segment max 1000 meters
 - (PHY – 1700 meters)
 - Spur segment max 200 m
- PHY and (SPI) MAC-PHY components available

Currently under IEEE 802.3 standardisation




Figure 5: Example topology for long cable reach with up to 1000 m between switches on the trunk

■ Future: Building Automation?

<https://encrypted-tbn1.gstatic.com/images?q=tbn:ANd9GcTIGTUf7MiR4D3CM7UYeeundXHwVpoXWGL6T4YB0HdD-bd5nrEDQ>

https://static1.squarespace.com/static/5c33ae96f407b4bfc5884978/t/60c7ef3a4807d14ca771825d/1623715647354/Ethernet-APL_Ethernet-To-The-Field_EN_FINAL_June-2021.pdf

Reference: https://www.ieee802.org/3/1TPCESG/public/BroadR_Reach_Automotive_Spec_V3.2.pdf

The MAC (2)

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack


Three ways of getting frame to memory

1.) direct collect by CPU (if buffers in MAC)

2.) MAC initialises DMA transfer – both types of MAC


3.) MAC writes directly into memory

The entire system is affected in different ways by the decision of MAC architecture



The MAC (3)

- Media Access Controller – two general architectures
 - Integrated buffer
 - External buffer



Learning Aim: The student will be able to explain the difference between buffering architectures and their effect on data transfer in the MAC

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication**
- Physical Layer
- Session 3: Communication Stacks**
- Communication Stack

The MAC (4)

- Session 1: Constrained Devices
 - Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
 - Physical Layer
- Session 3: Communication Stacks
 - Communication Stack

Buffer Descriptors

An array of descriptors with pointer to memory and a status word

Status word includes any errors, frame size and typically a «continue» bit


If, f.i. bit set then after frame tx or rx use the next buffer descriptor / next buffer descriptor is valid

Data can be transferred word for word f.i. using DMA

Pointer to memory	
Frame status	1
Pointer to memory	
Frame status	1
Pointer to memory	
Frame status	0
Pointer to memory	
Frame status	1

Learning Aim: The student will be able to program a buffer descriptor

IoT Lecture 7 – Communication Stacks




https://en.wikipedia.org/wiki/Protocol_stack


Data Buffering Memory (1)

- Session 1: Constrained Devices
 - Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
 - Physical Layer
- Session 3: Communication Stacks
 - Communication Stack

Simple HW and Software Systems Buffering: Single Buffer



One of these gets dropped



Example Schedule Single Buffer


**Processor synchronous with
communication -> Critical Activities**

Learning Aim: The student will be able to explain the three data buffering methods, single, double and triple


Data Buffering Memory (2)

- Session 1: Constrained Devices
 - Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
 - Physical Layer
- Session 3: Communication Stacks**
 - Communication Stack

Simple HW and Software Systems Buffering: Double Buffer



Will get into trouble later than sooner



Example Schedule Double Buffer

Data Buffering Memory (3)


Simple HW and Software Systems Buffering: Triple Buffer

Session 1: Constrained Devices


- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer

Session 3: Communication Stacks

- Communication Stack



Triple buffer completely decouples CPU from communication



Example Schedule Triple Buffer

The Software (1)

Sequence diagram describes the passage of an IP frame through a PROFINET node under eCOS using BSD2000 stack. Application is a simple UDP server

Session 1: Constrained Devices


- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack



Naïve Stack Architecture

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication


- Physical Layer

Session 3: Communication Stacks

- Communication Stack

**Each layer
maintains its own
buffers**

Standard OSI-Compatible
Stack Data Flow



Learning Aim: The student will be able to explain the two presented stack architectures

The Copy Problem

Sequence diagram describes the passage of an IP frame through a PROFINET node

memcpy() !!!


- ## Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational

- Session 2: Ethernet Communication
 - Physical Layer

Session 3: Communication Stack

 - Communication Stack



Zero Copy Stack

Zero Copy Stack

Session 1: Constrained Devices


- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack



Learning Aim: The student will be able to explain zero-copy stack

Scatter-gather

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power


Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

Scatter Gather Mechanism



Learning Aim: The student will be able to explain scatter-gather DMA

Learning Aim: The student will understand that an optimised stack architecture utilises scatter-gather and zero-copy features

Scatter-gather – example code

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power
- Session 2: Ethernet Communication
- Physical Layer

Session 3: Communication Stacks

- Communication Stack

```

static void sopc_bulk_aio_2p_recv( struct eth_drv_sc *sc, struct eth_drv_sg *sg_list, int sg_len) {
    cyg_uint32 pos;
    cyg_uint8 i;

    sopc_bulk_aio_2p_priv_data *cpd = (sopc_bulk_aio_2p_priv_data *)sc->driver_private;
    debug1_printf("recv packet with elem %d len %d\n", sg_len, cpd->recv_len);


    pos = 0; |
```

```

        for (i=0; i<sg_len; i++) {
            debug1_printf("r: sg_list elem %d, len %d\n", i, sg_list[i].len);
            memcpy((void*)sg_list[i].buf, &(cpd->recv_buffer[pos]), sg_list[i].len);
            pos+=sg_list[i].len;
        }
    }
```

There's memcpy again !!

lucky we have HW guys



Linux

Session 1: Constrained Devices


- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication


- Physical Layer

Session 3: Communication Stacks

- Communication Stack



- The NIC maintains a ring buffer of descriptors to **sk_buffs**. These are initialised by the kernel and point to kernel locations where frames and their meta data can be stored. The packet data is part of this structure and has pointers to the data in the structure.
- When a packet is passed from IP to TCP then the **sk_buff** is cloned. The metadata is copied but the packet data isn't. The clone is then passed from IP to TCP - the packet itself remains in kernel memory



The student will be able to explain how Linux Implements zero-copy

Learning Aim: The student will be able to explain how Linux implements zero-copy

Linux – sk_buff

The process for transmit is broadly similar
below the first sk_buff entries

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer


Session 3: Communication Stacks

- Communication Stack

```

231 struct sk_buff {
232     /* These two members must be first. */
233     struct sk_buff *next;
234     struct sk_buff *prev;
235
236     struct sock      *sk;      /* owner socket */
237     struct skb_timeval tstamp; /* arrival time */
238     struct net_device *dev;    /* output dev */
239     struct net_device *input_dev; /* input dev */
240
241     union {
242         struct tcphdr *th;
243         struct udphdr *uh;
244         struct icmphdr *icmph;
245         struct igmphdr *igmph;
246         struct iphdr   *ipiph;
247         struct ipv6hdr *ipv6h;
248         unsigned char *raw;
249     } h;      // <--- Transport header address
250
251     union {
252         struct iphdr   *iph;
253         struct ipv6hdr *ipv6h;
254         struct arphdr  *arph;
255         unsigned char *raw;
256     } nh;      // <--- Network header address
257
258     union {
259         unsigned char *raw;
260     } mac;    // <--- MAC header address

```



<http://www.ece.virginia.edu/cheetah/documents/papers/TCPlinux.pdf>

Linux – User Interface Sockets (1)

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

The socket is an interface methodology developed at UC Berkeley


It is optimised for streaming interfaces

It has become the de-facto standard networking interface

In our world we usually identify it with TCP and UDP protocols however it also supports other protocols such as AppleTalk and X.25

It also supports raw sockets – Layer 2/3 packets

Like most Linux I/O sockets are realised as file operations using SOCKFS and socket/close/read/write operations



Learning Aim: The student will be able to explain how Linux implements sockets

Learning Aim: The student will be able to explain how Linux implements client-server communications

Linux – Sockets (2)

■ Create a socket

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

```

/* Create socket */
socket_desc = socket( AF_INET , SOCK_STREAM , 0 );
if( socket_desc == -1 )
{
    fprintf( stderr, "could not create socket. exiting ...\\n" );
    exit( EXIT_FAILURE );
}

server.sin_addr.s_addr = inet_addr( "195.176.255.236" );
server.sin_family = AF_INET;
server.sin_port = htons( 80 );

/* Connect to remote server */
if( connect( socket_desc, ( struct sockaddr * )&server , sizeof( server ) ) < 0 )
{
    fprintf( stderr, "connect error. exiting ...\\n" );
    exit( EXIT_FAILURE );
}

/* Send some data (in this case a http request) */
message = "GET / HTTP/1.1\\r\\n\\r\\n";
if( send( socket_desc, message, strlen( message ), 0 ) < 0 )
{
    fprintf( stderr, "failed to send data. exiting ...\\n" );
    exit( EXIT_FAILURE );
}

```

IPV4 TCP

Could be write()

Connect to this address/port

Send HTTP GET

Linux – Sockets (3)

Session 1: Constrained Devices

- Definitions
- Benchmarking
- Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- Communication Stack

Wait for reply

```
/* receive a reply from the server */  
if( recv( socket_desc, server_reply , 2048 , 0 ) < 0 )  
{  
    fprintf( stderr, "failed to receive data from server. exiting ...\\n" );  
    exit( EXIT_FAILURE );  
}
```

Could be read()

Close socket – terminate TCP connection

```
errsv = close( socket_desc );  
simple_err_check( "failed to close socket file descriptor.", errsv );
```

Linux – Raw Sockets

Session 1: Constrained Devices

- Definitions
 - Benchmarking
 - Computational Power

Session 2: Ethernet Communication

- Physical Layer

Session 3: Communication Stacks

- **Communication Stack**

Use IPV4 infrastructure

```
/* Submit request for a socket descriptor to look up interface. */
if( ( sd = socket( AF_INET, SOCK_RAW, IPPROTO_RAW ) ) < 0 )
{
    fprintf( stderr, "socket() failed to get socket descriptor for usir
    print_usage();
    exit( EXIT_FAILURE );
}
```

The diagram shows two annotations with red arrows pointing to specific parts of the code. One arrow points to the line 'sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW)' with the label 'raw socket' below it. Another arrow points to the line 'IPPROTO_RAW' with the label 'No IP header' below it.

Raw Sockets allow you build your own protocols on user level

```
/* Sender hardware address (48 bits): MAC address */
memcpy( &arphdr.sender_mac, src_mac, 6 * sizeof( uint8_t ) );
```